

Production Setup Guide

End-to-end installation, configuration and verification for the GeoSphere 360 WebGIS dashboard and on-premises survey pipeline

DOCUMENT	GeoSphere 360 — Production Setup Guide
REVISION	1.0
APPLIES TO	geosphere-360-dashboard
AUDIENCE	Client IT / infrastructure team, and client survey operations leads
SCOPE	Cloud deployment, database, on-premises services, configuration, verification
STATUS	For implementation

Read this first. This platform performs **no server-side image processing** and requires **no GPU**. Panorama quality analysis runs in the operator's browser on a WebGL shader path; image editing is performed by operators in desktop software on four workstation PCs, which this platform observes. Do not procure GPU hardware or CUDA toolchains for this system. See **Section 1.2**.

AT A GLANCE

GPU REQUIRED No	TUNNEL HOSTNAMES 5	DATABASE MIGRATIONS 28 files
ON-PREM SERVICES 5	VERIFICATION TIERS 2	HOSTING TARGET Cloudflare Pages

Known limitations are catalogued in **Section 14**. Two require a client decision before go-live: panorama bucket privacy (4.1) and database-enforced deletion approval (6.3).

How to read this guide, and what it assumes.

Purpose

This guide takes a client from an empty account to a verified, running GeoSphere 360 installation. It is written to be followed in order, by someone who has not seen the system before. Every step states what to run, the expected result, and what to do when it differs.

How the guide is organised

The work divides into four phases. Each phase can be handed to a different engineer, and later phases depend on earlier ones.




Figure 1. The four install phases and their dependencies.

Conventions

CONVENTION	MEANING
<ANGLE_BRACKETS>	A placeholder. Replace the text <i>and</i> the brackets.
VITE_SOMETHING	An environment variable name. Names are case-sensitive.
Windows / Linux	PowerShell versus bash. Commands are otherwise identical.
Section <i>n</i>	A numbered section of this guide.
Blue callout	Information you will need.
Amber callout	A common mistake. Check here when something does not match.
Red callout	A known limitation or blocker. Read before go-live.

**Build-time versus runtime configuration**

Variables prefixed `VITE_` are compiled into the browser bundle at build time. Adding or changing one requires a **rebuild** — editing the variable and restarting the application does nothing. Variables without that prefix are read server-side by Cloudflare and never reach the browser. Sections 8.3 and 8.4 list which is which.

- 1 **About this system**
What it is, what it is not, architecture, components
- 2 **Prerequisites**
Accounts, domain and hostnames, hardware, software versions
- 3 **Database**
Project creation, PostGIS, the migration sequence, verification
- 4 **Storage buckets**
Panorama bucket, geometry bucket, privacy implications
- 5 **First administrator**
Manual bootstrap, role resolution, adding further users
- 6 **Roles and row-level security**
The four roles, what the database actually enforces
- 7 **Cloudflare Tunnel**
Named tunnels, configuration, DNS routing, running as a service
- 8 **Cloudflare Pages deployment**
Build settings, build-time and runtime variables, hosting constraint
- 9 **NAS layout** Required folder tree, depth rules, subgrid naming
- 10 **NAS worker**
Install, configure, run, verify
- 11 **Station agents**
One per workstation, configuration, firewall, autostart
- 12 **First-run configuration**
Project, settings, station address setup, roles
- 13 **Verification**
Tier A and Tier B acceptance, smoke test, sign-off sheet
- 14 **Known limitations**
Product defects that affect a client installation
- 15 **Troubleshooting**
Symptom to cause, including silent failures
- 16 **Reference**
Environment variables, migrations, HTTP surfaces, roles

What GeoSphere 360 is, how it is put together, and what it does not do.

1.1 What it is

GeoSphere 360 is a browser-based WebGIS dashboard for 360 degree mobile-mapping survey data, together with a small set of on-premises services that connect it to a survey NAS and to four photogrammetry workstations.

A surveyor captures 360 imagery and GNSS/IMU telemetry on a vehicle rig. Operator PCs run desktop photogrammetry software to blur, stitch, enhance and retouch the imagery. The dashboard coordinates that work: it shows which subgrid has been captured, what stage each workstation has reached, which frames failed quality inspection, and what has been published to the WebGIS layer.

1.2 What it is not

This distinction drives hardware purchasing, so it is stated plainly and early.



No server-side image processing

The platform does **not** run a server-side image pipeline. There is no GPU worker, no batch enhancement queue, and no image-processing job store.

- **Panorama quality analysis** — sharpness, privacy-blur verification, obstruction and glare detection — runs **in the operator's browser** on a WebGL shader path, with a CPU fallback.
- **Image editing** — blurring, stitching, enhancement, nadir removal — is performed by operators in desktop software on the four workstations. The dashboard observes their progress; it does not drive them.

You do not need: an NVIDIA GPU, CUDA, PyTorch, OpenCV, a GPU workstation, or a GPU queue service. Standard office hardware is sufficient. Section 2.3 gives the actual list.

1.3 Architecture

The system splits into a cloud half, which the client's users reach over HTTPS, and an on-premises half that never touches the public internet. A Cloudflare Tunnel bridges them.




Figure 2. System architecture. The browser never receives an on-premises credential — Pages Functions validate the session and inject server-side secrets before proxying. The NAS is reachable only by the worker and the agents.

COMPONENT	RESPONSIBILITY	AUTHORITATIVE FOR
Cloudflare Pages	Serves the web application; hosts the private-API proxy	—
Pages Functions	Validate the signed-in session, then proxy to on-premises services using server-side credentials	Never exposes an on-premises credential to the browser
Supabase	Spatial database, authentication, object storage	All persistent state
NAS worker	Read-only view of the survey filesystem and storage telemetry	What is on disk
Station agents x4	Report workstation progress; perform NAS rename and bucket upload	Per-workstation work state
Operator browser	Panorama quality analysis on a WebGL path	Quality verdicts



Design principle

The database is the single source of truth. The NAS is treated as read-only by the dashboard: the worker rejects any path that resolves outside its configured base, so a malformed request cannot reach unrelated files on the client network.

Accounts, domain, hostnames, hardware and software — everything to acquire before starting.

2.1 Install order at a glance

Perform these in sequence. Steps 6 to 9 are the on-premises half and may be carried out in parallel with steps 1 to 5 by a different engineer.

#	ACTION	SECTION	ESTIMATE
1	Create Cloudflare and Supabase accounts; attach a domain	2.2, 2.4	1 day
2	Create the Supabase project; enable PostGIS first	3.1, 3.2	30 min
3	Apply migrations <code>0001</code> – <code>0010</code> then <code>0012</code> – <code>0029</code>	3.3	45 min
4	Create the <code>MMS_PIC</code> storage bucket (public)	4.1	5 min
5	Create the first administrator	5	15 min
6	Build the NAS folder tree	9	1 day
7	Install and run the NAS worker	10	30 min
8	Install and run four station agents	11	2 hours
9	Create five Cloudflare Tunnels	7	1 hour
10	Deploy to Cloudflare Pages	8	30 min
11	Configure the application	12	45 min
12	Verify the installation	13	30 min

2.2 Accounts and services

NEED	WHY	NOTES
Cloudflare account	Hosts the web deployment, the tunnels and optional access policies	A free account is sufficient to start
A domain on Cloudflare DNS	The tunnel hostnames must be real, resolvable HTTPS names	Section 2.4 — you need five of them
Supabase project	PostgreSQL and PostGIS, authentication, object storage	Supabase Cloud, or self-hosted Docker
Photogrammetry software licences	Runs on the four workstations	Not part of this platform
A WebGIS application URL	The embedded interactive map	Required at build time; there is no fallback

2.3 Hardware

No GPU is required for any component of this platform (Section 1.2).

NAS / worker PC	4 cores, 8 GB RAM, 100 GB free disk	8 cores, 16 GB RAM, SSD
Each station PC (x4)	4 cores, 8 GB RAM, SSD	8 cores, 16 GB RAM, SSD
Network	1 GbE between the PCs and the NAS	10 GbE for large surveys
NAS capacity	Raw + stitched + final + deliverables	Plan roughly three times the raw survey volume
Cloud	Managed Supabase, free Cloudflare tier	Scale the database for archive volumes

 **Size the workstations for the photogrammetry software**
The station PCs also run stitching and image-editing software, which is far more demanding than this platform. Choose their specification for that workload, not for the dashboard.

2.4 Hostnames — plan for five

The dashboard reaches the on-premises services through five separate tunnels. This is the single most commonly under-provisioned item, so plan for it before ordering certificates or DNS.

ONE DNS ZONE, FIVE TUNNEL HOSTNAMES

nas-api.example.com
NAS worker · required for NAS features

pc1-agent
station · blur

pc2-agent
station · stitch

pc3-agent
station · lightroom

pc4-agent
station · photoshop

All five must be **HTTPS**. The station-agent proxy rejects any plain-http origin outright.

Free *.trycloudflare.com hostnames are **not usable** — they change on every restart, so agents would drop constantly.

The web application itself may first be deployed on the free *.pages.dev hostname before a domain is attached.

Optional sixth hostname: pc2-remote for an embedded remote desktop (Section 8.7).

Figure 3. Required hostname allocation.

2.5 Software versions

COMPONENT	VERSION	NOTES
Node.js	18.x or 20.x LTS	Only needed on the machine that builds the application
npm	9.0 or higher	Ships with Node 18 and 20
Python	3.10 or higher	Runs the NAS worker and the station agents
cloudflared	current	Runs on the worker PC and each station PC
PostgreSQL	15 with PostGIS 3.3	Provided by Supabase

Project creation, PostGIS, the migration sequence and verification.

3.1 Create the project

Supabase Cloud — create a project in your organisation. **Self-hosted** — run the Supabase Docker stack. Record the project URL and the anon / publishable key; you will enter them in both the browser-facing and the server-facing configuration.

3.2 Enable PostGIS before the first migration


 **Ordering trap — enable PostGIS first**

The first migration creates a summary view **only if PostGIS already exists**. If it is absent at that moment the view is silently skipped, and PostGIS is only installed later by migration `0012`.

Do this first: Supabase → **Database** → **Extensions** → `postgis` → **Enable**. Then apply the migrations.

3.3 Apply the migrations

This project is not wired to a migration runner: there is no Supabase configuration file and no seed script, so `supabase db push` will not work. The migrations are applied as a hand-run, ordered sequence — 29 of the 30 files.



0011_security_tests.sql		Not a migration — verification only (Section 6.2)
11	0012_core_tables_and_rls.sql	PostGIS, panoramas, staging_panoramas, panoramas_view
12	0013_prune_blobbloated_metadata.sql	Strips oversized road-analysis state from user metadata
13	0014_road_analysis_state_rpc.sql	sec.save_road_analysis_state()
14	0015_projects.sql	projects registry
15	0016_per_project_isolation.sql	project_id across 12 tables, projects_adapt_legacy_rows
16	0017_project_delete_cascade.sql	projects_delete_cascade()
17	0018_map_shares.sql	map_shares, map_share_touch()
18	0019_map_shares_public_access.sql	Public read policy for shares
19	0020_datasets_operator_delete.sql	Operator delete rights on datasets
20	0021_road_analysis_geometry_storage.sql	Creates the private road-analysis-geometry bucket
21	0022_production_runs.sql	production_runs, attempts, releases, release files
22	0023_station_board_items.sql	station_board_items
23	0024_stage_event_ledger.sql	stage_event_ledger
24	0025_station_board_metric_unit.sql	Capture-point units on the station board
25	0026_hub_session_state.sql	hub_session_state
26	0027_reconcile_panorama_columns.sql	Latitude, longitude and heading columns, with backfill
27	0028_sync_file_inventory.sql	Storage-to-inventory triggers and backfill
28	0029_csv_file_name.sql	CSV filename column and backfill
29	0030_role_precedence.sql	Directory row becomes authoritative for role and status — makes Disable revoke access. Back-fills first, so it cannot lock anyone out
30	0031_rls_completeness_backstop.sql	Closes the two RLS gaps a fresh install leaves: survey_recycle_bin never had RLS enabled, and batch_logs had RLS but no policies (deny-all)



File name check

The file at position 18 is 0019_map_shares_public_access.sql — it grants **public** read access to share rows. Older notes refer to 0019_map_shares_auth_only.sql, which does not exist.

Ordering dependencies

The listed order already satisfies these, but they explain the failures you will see if the order is broken. **One of them is a known defect in this set**, covered in the warning immediately below the table.

CONSTRAINT	REASON
------------	--------

0018 before 0019	0018 adds <code>project_id</code> columns and a helper function
0018 before 0019	0019 re-grants and re-creates the policy
0023 before 0025	0025 warns and exits if the station board table is absent
0007 and 0016 before 0031	0031 backstops <code>survey_recycle_bin</code> and <code>batch_logs</code> , which those two create
Known defect: 0012 before 0004	0004 opens with <code>ALTER TABLE public.panoramas</code> , but <code>panoramas</code> is created by 0012 — eight files later, so on an empty database 0004 fails on its first statement



On a fresh install, do not stop at the first error

Because 0004 needs a table that 0012 creates, a fresh database fails at 0004 and the failure cascades: 0007 needs `deletion_requests`, and 0010 and 0030 need `user_accounts`.

If you stop there, 0030 never runs and Disable does not actually revoke access — accounts look revoked while their token claim still grants everything.

The Supabase SQL editor is safe: it runs statements independently, so 0012 can still create the tables 0004 wanted. With `psql`, omit `ON_ERROR_STOP=1`. 0031 then repairs the two RLS gaps this ordering leaves behind.

How to run them

Supabase SQL editor — paste and run each file in order, one at a time. Safe on a fresh install.

psql — note this deliberately does **not** fail fast, for the reason above:

```
# bash
export PGHOST=<db-host> PGDATABASE=postgres PGUSER=postgres PGPASSWORD=<password>
for f in supabase/migrations/00*.sql; do
  case "$f" in *0011_security_tests.sql) continue ;; esac
  # deliberately NOT ON_ERROR_STOP: 0004 needs a table 0012 creates, so it fails on a fresh db
  psql -q -f "$f" || echo " (continued past an error in $f"
done
```

The migrations are written to be idempotent, so re-running a file that partially applied is safe. After the full sequence, run the verification in 3.4.

```

-- security functions present (expect: can, get_app_role, is_role, normalize_role)
select proname from pg_proc p join pg_namespace n on n.oid = p.pronamespace
  where n.nspname = 'sec' order by 1;

-- PostGIS enabled
select extname from pg_extension where extname = 'postgis';

-- row-level security actually enabled on the privileged tables (expect several rows)
select relname, relrowsecurity from pg_class
  where relnamespace = 'public'::regnamespace and relrowsecurity
  order by 1;

-- the views the application reads
select table_name from information_schema.views
  where table_schema = 'public'
     and table_name in ('panoramas_view', 'panoramas_subgrid_summary');

-- GATE: directory-authoritative role resolver in place (expect NEW)
-- OLD means 0030 did not apply, and Disable does not revoke access
select case when prosrc like '%not Active%'
  then 'NEW' else 'OLD' end as resolver
  from pg_proc where proname = 'get_app_role';

-- RLS gaps: expect ZERO rows in both (spatial_ref_sys is PostGIS's own)
select c.relname as rls_missing
  from pg_class c join pg_namespace n on n.oid = c.relnamespace
  where n.nspname = 'public' and c.relkind = 'r'
     and not c.relrowsecurity and c.relname <> 'spatial_ref_sys';

-- RLS on but ZERO policies = deny-all. Expect ZERO rows.
select c.relname as no_policies
  from pg_class c join pg_namespace n on n.oid = c.relnamespace
  where n.nspname = 'public' and c.relkind = 'r' and c.relrowsecurity
     and not exists (select 1 from pg_policies p
  where p.schemaname = 'public' and p.tablename = c.relname);

```

If `panoramas_subgrid_summary` is missing, PostGIS was not installed when migration `0001` ran. Enable PostGIS and re-run the view block from `0001` (Section 3.2).

If the resolver reads **OLD**, the sequence stopped before `0030` — most often because `psql` was run with `ON_ERROR_STOP=1` and stopped at `0004`. Re-run the sequence tolerantly, then `0030`. If either RLS query returns rows, apply `0031`.

Which buckets to create, and the privacy decision you must make consciously.

4.1 Panorama bucket — create it manually, and make it public

No migration creates the panorama bucket. Create it yourself: Supabase → **Storage** → **New bucket**.

SETTING	VALUE
Name	MMS_PIC
Public bucket	ON



Why it must be public

The dashboard builds panorama URLs of the form

`https://<project>.supabase.co/storage/v1/object/public/MMS_PIC/<file>` and hands them directly to image and tile loaders. There is no signed-URL code path for panoramas. With a private bucket, every panorama renders broken.



Conflicting internal documentation

Some internal notes describe this bucket as private. Making it private is only viable if you *also* add a storage read policy **and** have the application switched to signed URLs. The second half is **not implemented**. Treat "public" as the only supported configuration today.



Privacy decision — make it consciously

With a public bucket, the object path is the only protection on an image, and survey filenames typically encode the subgrid and frame index. If your imagery is confidential, agree an alternative with your integrator **before go-live**: either keep raw imagery off this bucket and serve previews through the NAS worker, or request the signed-URL capability as a development item.

4.2 Road-analysis geometry bucket — automatic

Migration 0021 creates a bucket named `road-analysis-geometry` with private access and authenticated-only storage policies. **No action is required**. If it is missing, Road Analysis geometry saves fall back to browser-local storage, so nothing breaks visibly — check for it if layer saves appear not to persist.

4.3 A bucket you do not need

Some older documentation instructs you to create a `vector_layers` bucket. The current code never reads or writes it. Skip it.

4.4 Bucket summary

BUCKET	HOW IT IS CREATED	PRIVACY
MMS_PIC	Manually, in the dashboard	Public
road-analysis-geometry	Migration 0021	Private

Bootstrapping the first account, which cannot be done from inside the application.

5.1 Why this must be done by hand

There is no seed script, no migration-runner configuration and no sign-up path in the application. Meanwhile the user-management screen requires the `manageUsers` capability, which only an Administrator holds. The first Administrator therefore **cannot** be created from within the product — you must seed them.




Figure 5. First-administrator bootstrap. All four steps are required, and step 3 must complete **before** the account's first sign-in.

5.2 The four steps

1 Create the authentication user

In the Supabase dashboard go to **Authentication** → **Users** → **Add user**. Enter an email address and password, or send an invitation. Note the address — it is used in step 3.

2 Set the role metadata

Open the new user and edit their metadata. Set **App Metadata** first, because that is what lands in the JSON Web Token and is the fastest path to a working administrator.

```
{ "role": "Administrator" }
```

Set **User Metadata** to the same value, because the dashboard reads User Metadata before App Metadata.

3 Insert the directory row

```
insert into public.user_accounts (name, email, role, status)
values ('<Full Name>', '<email>', 'Administrator', 'Active')
on conflict (email) do update
  set role = 'Administrator', status = 'Active';
```

4 Verify before signing in

```
select sec.get_app_role(); -- must return: Administrator
```

If it returns `Viewer`, one of steps 2 or 3 is missing or incomplete.



Complete step 3 before the first sign-in

On first sign-in the application writes its own `user_accounts` row using the role it resolved from metadata, **defaulting to Viewer**. Running step 3 afterwards corrects it, but setting only the metadata afterwards does not, because the cached Viewer row already exists.

two independent systems must agree, and neither reads the other.

SYSTEM	READS FROM	CONSEQUENCE
Dashboard interface	Authentication metadata on the signed-in user	Decides which controls are shown
Database	The token's role claim, then the <code>user_accounts</code> row	Decides which rows are visible

The dashboard tries, in order: `user_metadata.role`, `raw_user_meta_data.role`, `user.role`, `app_metadata.role`, `raw_app_meta_data.role`. For a real account, `user.role` holds the string `"authenticated"`, which is not an application role and normalises to **Viewer**. If no metadata key is set, the result is Viewer with no warning shown.

5.4 Adding the remaining users

Authentication lives entirely in Supabase. **Administration** → **Users does not create** accounts — it manages the role and access of people who already exist in Supabase.

1 Create the account in Supabase

In the Supabase dashboard go to **Authentication** → **Users** → **Add user**, with their email address, a password, and **Auto Confirm User** enabled.

2 Set the role as App Metadata

On that user, set **App Metadata** to the role they need:

```
{ "role": "Survey Operator" }
```

Valid values are `Administrator`, `Survey Operator`, `QA Inspector` and `Viewer`. Optionally set **User Metadata** to the same value, because the interface reads User Metadata first.

3 They appear on first sign-in

The user shows up in **Administration** → **Users** automatically the first time they sign in — that is when their directory row is written. Before then there is no row for them. From that point the dashboard manages their role and access.

ⓘ Two things that catch people out

- **Set App Metadata, not only User Metadata.** The database resolves the role from the token's role claim, which Supabase sources from App Metadata. User Metadata alone makes the interface look correct while the database still sees `Viewer` — the user sees controls they cannot actually use.
- **Disable revokes permissions, not the account.** It sets the directory row to `Disabled`, which migration `0030` resolves to read-only. The Supabase sign-in account still exists. To stop someone signing in at all, remove or disable the user in Supabase.

Grant individual capabilities in **Administration** → **Roles**.

The four roles, and — importantly — which controls the database actually enforces.

6.1 The four roles

ROLE	INTENDED FOR
Administrator	System owners. Full access, including user management and system settings.
Survey Operator	Field and production staff. Intake, pipeline operations, dataset management.
QA Inspector	Quality staff. Quality inspection and sign-off.
Viewer	Stakeholders. Read-only, plus report export.

6.2 Verifying the security boundary

Tier 1 — always runnable

Run the Section 3.4 queries, then paste and run all of `supabase/migrations/0011_security_tests.sql` . Its first, second and final sections must complete **without raising an exception**.

 **The SQL editor cannot prove the boundary**
The Supabase SQL editor runs as the `postgres` superuser, which **bypasses row-level security entirely**. Tier 1 proves the migrations applied and that role resolution is wired. It cannot prove that a Viewer is denied.

Tier 2 — a real boundary test

Requires a Viewer test account and a direct database connection:

```
psql "$SUPABASE_DB_URL" -v ON_ERROR_STOP=1 <<'SQL'
BEGIN;
SET LOCAL ROLE authenticated;
SET LOCAL request.jwt.claims = '{"role":"authenticated","email":"viewer@test.local"}';
select sec.get_app_role(); -- must print: Viewer
insert into public.project_settings (id, settings)
  values ('rls_probe', '{}':jsonb);
-- EXPECT: ERROR new row violates row-level security policy    => PASS
ROLLBACK;
SQL
```

The insert **must fail** with a row-level-security error. If it succeeds, the boundary is broken — stop and resolve before go-live. Then repeat with an Administrator token and the same statement: that one **must succeed**.

6.3 What is actually enforced

 **Eight of twenty-one capabilities are enforced by the database**
The remaining controls gate the **interface only** — the button is hidden or disabled, but a direct API call is not blocked. If your compliance requirements depend on server-side enforcement, treat the other thirteen as advisory and restrict roles accordingly.

ENFORCED BY ROW-LEVEL SECURITY	ADMINISTRATOR	SURVEY OPERATOR	QA / VIEWER
--------------------------------	---------------	-----------------	-------------

Capability	Survey Operator	Administrator	Survey Manager
<code>deleteData</code>	Yes	Yes	No
<code>manageUsers</code>	Yes	No	No
<code>manageSettings</code>	Yes	No	No
<code>manageDatasets</code>	Yes	Interface only	Interface only
<code>reviewQaqc</code>	Yes	Interface only	Interface only
<code>approveDeletions</code>	Yes	No	No

**Deletion approval is an interface guard**

The `deleteData` capability is granted at the database level to **both** Administrator and Survey Operator. The "require administrator approval for delete" setting only gates the user interface. To enforce approval, do not grant the Survey Operator role where your process requires it.

Exposing the on-premises services to the dashboard without opening any port to the internet.

The dashboard is served over HTTPS from Cloudflare's edge. Your on-premises services sit on a private network.

Cloudflare Tunnel bridges the two: `cloudflared` runs on each machine and makes **outbound-only** connections to Cloudflare.



No inbound firewall rule needed for the tunnel

Because `cloudflared` dials outward, nothing needs to accept an inbound connection. You may still restrict port 8000 at the firewall so that the services remain reachable only from the local network (Section 11.5).

7.1 Install cloudflared

```
# Windows
winget install --id Cloudflare.cloudflared

# Linux / macOS - see the Cloudflare downloads page for the current package
# https://developers.cloudflare.com/cloudflare-one/connections/connect-networks/downloads/

cloudflared --version
```

7.2 Create one tunnel per machine

Create a tunnel **on each machine that hosts a service**: one on the worker PC, and one on each of the four station PCs. Five in total.

```
# on the worker PC
cloudflared tunnel login
cloudflared tunnel create nas-worker

# on each station PC, using a distinct name
cloudflared tunnel create pc1-agent
cloudflared tunnel create pc2-agent
cloudflared tunnel create pc3-agent
cloudflared tunnel create pc4-agent
```

`login` opens a browser to authorise your Cloudflare account. `create` prints a **tunnel identifier** and writes a credentials file under `~/.cloudflared/<TUNNEL-ID>.json`. **Keep that file** — it is the tunnel's identity, and the next step references it.

7.3 Write the ingress configuration

Each tunnel needs a configuration file mapping public hostnames to local services. Create

`%USERPROFILE%\cloudflared\config.yml` on Windows, or `~/.cloudflared/config.yml` elsewhere.

Worker PC

```
tunnel: <TUNNEL-ID>
credentials-file: C:\Users\<USER>\.cloudflared\<TUNNEL-ID>.json

ingress:
  - hostname: nas-api.example.com
    service: http://localhost:8000
```

```

tunnel: <TUNNEL-ID>
credentials-file: C:\Users\<USER>\.cloudflared\<TUNNEL-ID>.json

ingress:
  - hostname: pc1-agent.example.com
    service: http://localhost:8000
  - service: http_status:404

```

ⓘ The trailing catch-all rule is required

`cloudflared` refuses to start unless the final ingress rule is a catch-all. Keep the `- service: http_status:404` line exactly as shown.

If you also want an embedded remote desktop for a station, add a second hostname pointing at your remote-desktop service (commonly `http://localhost:5901`) — see Section 8.7.

7.4 Point DNS at the tunnel

```

cloudflared tunnel route dns nas-worker nas-api.example.com
cloudflared tunnel route dns pc1-agent pc1-agent.example.com
cloudflared tunnel route dns pc2-agent pc2-agent.example.com
cloudflared tunnel route dns pc3-agent pc3-agent.example.com
cloudflared tunnel route dns pc4-agent pc4-agent.example.com

```

This creates the DNS records automatically.

7.5 Run it as a service

Both options start the tunnel with the machine.

```

# Windows - registers a service that starts at boot
cloudflared service install

# then set the service to start automatically in services.msc

```

```

# Linux - systemd unit at /etc/systemd/system/cloudflared.service
[Unit]
Description=cloudflared
After=network-online.target

[Service]
Type=simple
ExecStart=/usr/local/bin/cloudflared --no-autoupdate run --token <TUNNEL-TOKEN>
Restart=on-failure

[Install]
WantedBy=multi-user.target

```

```
sudo systemctl enable --now cloudflared
```

The single-line **token** form is the simplest: in the Cloudflare Zero Trust dashboard, open the tunnel, copy the install command, and run it.

```
curl https://nas-api.example.com/health
```

An **HTTP 401** is the **correct answer** when no token is supplied — it proves the tunnel is routing and the worker is enforcing its guard:

```
{ "detail": "Unauthorized." }
```

Now with the token, which should return the live response:

```
curl -H "Authorization: Bearer <NAS_WORKER_TOKEN>" https://nas-api.example.com/health  
  
{ "status": "ok", "nas_base": "X:\\nas\\360_images", "worker": "NAS-PC-01" }
```

7.7 Optional — Cloudflare Access

For an additional gate at the origin, create a Cloudflare Access service token and set `CF_ACCESS_CLIENT_ID` and `CF_ACCESS_CLIENT_SECRET` (Section 8.4). The Pages Functions forward them automatically. This is defence in depth *on top of* the session check, not a replacement for it.

Build settings, the hosting constraint, and every variable to set.

8.1 Hosting constraint — read this first



Static-only hosting does not work with this platform

The private API — NAS scans, image previews, station probes, rename and upload — is implemented using the **Cloudflare Pages Functions** runtime: each handler exports an `onRequest` function and reads secrets from `context.env`.

Serving the built folder from nginx, Caddy or plain object storage **does not execute those handlers**. Under a static-only host:

- every private API route returns the application's HTML instead of JSON,
- the interface degrades **silently**, reporting "unavailable" rather than raising an error,
- and the session check that guards those routes does not run at all.

Deploy to Cloudflare Pages. A Pages-compatible edge runtime also works. If you must self-host, you need to build and maintain your own Pages-compatible server; that component does not ship with this product.

8.2 Create the project

Connect the repository in the Cloudflare dashboard:

SETTING	VALUE
Project name	Your choice
Root directory	Repository root
Build command	<code>npm run build</code>
Build output directory	<code>dist</code>
Production branch	<code>main</code> (or your deploy branch)



Root directory must be the repository root

Cloudflare detects the `functions/` directory there and deploys it automatically. If the root is set to `dist`, all seven function handlers are **silently omitted** and you get a dashboard with no private API and no error message. This is the single most damaging misconfiguration in this section.

A redirect rule in the project's public assets handles single-page route fallback for you. No server rewrite configuration is needed.

8.3 Build-time variables

These are **compiled into the browser bundle**. Set them under **Settings** → **Environment variables**. Only three are required.

VARIABLE	REQUIRED	PURPOSE
<code>VITE_SUPABASE_URL</code>	Yes	Project base URL. No fallback; the application warns if absent

VITE_MAP_URL	Yes	WebOS application only. No code fallback — unset means every embedded map is a blank frame
VITE_WORKER_API_MODE	No	<code>proxy</code> (default) or <code>direct</code>
VITE_STATION_AGENT_MODE	No	<code>proxy</code> in production, <code>direct</code> in development
VITE_SUPABASE_BUCKET	No	Defaults to <code>MMS_PIC</code>
VITE_SENTRY_DSN	No	Error reporting. Omit to disable entirely

**Never put a secret in a build-time variable**

Anything prefixed `VITE_` is inlined into the shipped bundle and is readable by anyone who views the page source. A service-role key, the NAS worker token or a station-agent token placed here would be **public**. Server secrets belong in Section 8.4.

**Changing a build-time variable requires a rebuild**

If you add a variable **after** a deployment, the running site is unchanged. Trigger a new build — push a commit, or use retry in the Pages dashboard.

**A retired name that does nothing**

`VITE_NAS_API_ENABLED` appears in some older notes but is **not read by any code**. Setting it has no effect. Proxy behaviour is derived from `VITE_WORKER_API_MODE` and the presence of the runtime secrets in Section 8.4.

8.4 Runtime variables and secrets

Read **server-side** by the Pages Functions; never sent to the browser. Mark them **Secret**.

VARIABLE	REQUIRED	PURPOSE
SUPABASE_URL	Yes	Session validation target
SUPABASE_ANON_KEY	Yes	Session validation key
NAS_API_URL	Yes	Worker tunnel origin, no trailing slash
NAS_WORKER_TOKEN	Yes	Must match the worker's token exactly
NAS_IMAGE_TOKEN_SECRET	No	Signing key for short-lived preview URLs; falls back to the worker token
STATION_AGENT_URLS	For the board	JSON map of station id to HTTPS agent origin
STATION_AGENT_TOKENS	No	JSON map of station id to that agent's token
STATION_AGENT_TOKEN	No	Single fallback token
CF_ACCESS_CLIENT_ID	No	Cloudflare Access service token, forwarded upstream
CF_ACCESS_CLIENT_SECRET		

`STATION_AGENT_URLS` example — valid station ids are exactly `blur`, `stitch`, `lightroom` and `photoshop` :

```
"lightroom": "https://pc3-agent.example.com",
"photoshop": "https://pc4-agent.example.com" }
```

8.5 Deploy

Push to the connected branch. Cloudflare Pages builds automatically — there is **no deploy automation in this repository**, so the Git integration is the entire mechanism.

8.6 Local development

```
npm install
cp .env.example .env          # Windows PowerShell: copy .env.example .env
```

Set these so the development server proxies the on-premises API exactly as production does:

```
VITE_SUPABASE_URL=https://<project>.supabase.co
VITE_SUPABASE_ANON_KEY=<anon key>
VITE_MAP_URL=<your WebGIS URL>

NAS_API_URL=https://nas-api.example.com
NAS_WORKER_TOKEN=<same token as the worker>
```

The last two are deliberately **not** prefixed with `VITE_`. The development server reads them and injects the authorisation header, so the browser never receives the token. If both are absent, no proxy is registered and the interface reports "worker proxy is not configured" rather than silently guessing a NAS address.


```
npm run dev
```

8.7 Remote desktop (optional)

To embed a station's live desktop, run a browser-based remote desktop service on that PC, expose it through a tunnel hostname (for example `pc2-remote.example.com`), and set the station's remote desktop URL in the application's provider settings.

A private `http://192.168.x.x` address will not work — the dashboard is served over HTTPS and browsers block mixed content. Native remote-desktop file downloads for operators on the local network continue to work.

8.8 Request path



- 1 The browser attaches the signed-in session token. It holds **no** on-premises credential.
- 2 The function validates the session against Supabase and reads its secrets from the server environment.
- 3 The function attaches the worker token and forwards the request through the tunnel.
- 4 The worker checks its token, resolves the path, and rejects anything outside the configured base.

If the function cannot reach the worker — typically a missing `NAS_API_URL` — the interface reports "unavailable" rather than erroring.

Figure 6. How a private request reaches the NAS, and where each credential is applied.

The folder structure the services read. Getting this wrong fails quietly, not loudly.

The worker and the agents read the survey filesystem directly, against **hard-coded stage folder names**. If the layout does not match, the corresponding scans return empty results rather than reporting an error — so build it exactly as shown.



Figure 7. Required folder tree. Colours mark the depth rules: amber **Grid 1** (literal, with a space), violet **Project-OUT**.

9.1 The depth rules

Three different conventions are in play. Mixing them up fails silently.

STAGE FOLDER	NEEDS PROJECT-OUT?	NEEDS GRID 1?	USED BY
00_Raw_data	No	Yes	Worker + agent
01_Metadata	No	Yes	Worker
02_Blurring	No	No	Agent only
03_Stitching	Yes	Yes	Worker + agent
04_Lightroom	No	No	Agent only
05_Final	Yes	Yes	Worker
DELIVERABLES	No	No	Agent only

 **Three things that commonly go wrong**

- Grid 1 contains a space.** On Linux, confirm the mount preserves it.
- The run folder name is a join key.** It links stitched imagery to its telemetry, so the run folder under **03_Stitching** and the run folder under **01_Metadata** must be **identical**. If they differ, coordinates will not be paired to imagery.
- 02_ , 04_ and DELIVERABLES have no depth rule.** Their first-level directories are read directly as subgrids.

Use the **uppercase GIS form**, for example `N93E70` or `N01E037`. The platform applies three different subgrid patterns for three different purposes, and the uppercase GIS form is the only shape that satisfies all three. Lowercase works on the worker but is handled inconsistently by the agents.

9.3 Access rights

ACCESS	REQUIREMENT
Read the whole survey tree	Required — worker and all four agents
Write to <code>DELIVERABLES</code>	Required — bucket upload
Write to <code>panoramas/</code>	Only for the rename action

The read-only service that exposes the survey filesystem to the dashboard.

10.1 Install

```
# Windows PowerShell
cd <repository>\worker
py -3.10 -m venv .venv
.\.venv\Scripts\Activate.ps1
pip install -r requirements.txt
```

```
# Linux
cd <repository>/worker
python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
```

Dependencies are deliberately minimal — a web framework, a server and a configuration loader. There is no image or scientific processing stack to install, and no model weights to download.

10.2 Configure

```
copy .env.example .env # Windows PowerShell
```

```
# Root of the survey tree. Every path the dashboard sends resolves under this,
# and any path escaping it is rejected with HTTP 400.
NAS_BASE_PATH=X:\nas\360_images

# Shared secret. Required on any non-loopback bind. The Pages Functions and
# the development proxy inject this header server-side; the browser never sees it.
NAS_WORKER_TOKEN=<generate a long random string>

# Seconds a storage response is cached before the volume is re-probed.
STORAGE_CACHE_TTL=30

# 1 = structured JSON logs to stdout.
NAS_LOG_JSON=0

# Identifier reported on /health, so several workers behind one endpoint
# can be told apart. Defaults to the machine hostname.
NAS_WORKER_ID=
```

Generate a token with:

```
python -c "import secrets; print(secrets.token_urlsafe(48))"
```



If the token is blank, the worker accepts unauthenticated requests

Only safe on a loopback bind with the tunnel in front. Never leave `NAS_WORKER_TOKEN` empty on a machine reachable from the local network.

```
uvicorn app:app --host 127.0.0.1 --port 8000
```

Bind to **localhost only**. The Cloudflare Tunnel on the same machine terminates TLS and is the sole public entry point. Do not bind `0.0.0.0` unless you also restrict the port at the firewall.

Run it as a service for autostart, using the same pattern as the agents in Section 11.6.

10.4 Verify

```
curl http://127.0.0.1:8000/health
```

```
{ "status": "ok", "nas_base": "X:\\nas\\360_images", "worker": "NAS-PC-01" }
```

Then test an authenticated data route:

```
curl -H "Authorization: Bearer <NAS_WORKER_TOKEN>" \
  "http://127.0.0.1:8000/api/nas-scan?action=subgrids"
```

Expected — your subgrid appearing here confirms the folder tree is correct:

```
{
  "success": true,
  "basePath": "X:\\nas\\360_images",
  "detectedSubgrids": ["N93E70"],
  "subgrids": [{ "code": "N93E70", "existsInStitching": true, "label": "N93E70" }]
}
```

An empty subgrid list means the tree does not match

Walk up from `NAS_BASE_PATH` and confirm that `03_Stitching/Project-OUT/Grid 1/` exists, spelled exactly as shown — this is the most common cause.

Also check storage telemetry:

```
curl -H "Authorization: Bearer <NAS_WORKER_TOKEN>" http://127.0.0.1:8000/api/storage
```

10.5 Run the worker tests

```
cd worker
PYTHONPATH=. python -m pytest tests -q
```

This covers the conditional-request behaviour of image previews and the six scan actions. The configuration loader must be installed, or the HTTP tests skip silently.

One small service per workstation, reporting progress and performing local actions.

11.1 Station roles

IDENTIFIER	PC	SOFTWARE	OUTPUT STAGE (DEFAULT)	READS
blur	1	Privacy / blur tool	02_Blurring	00_Raw_data
stitch	2	Stitching software	03_Stitching	02_Blurring
lightroom	3	Lightroom Classic	04_Lightroom	03_Stitching
photoshop	4	Photoshop batch	05_Final	04_Lightroom

Only these four values are accepted. Any other value produces an empty stage and a watch error.

11.2 Install

On **each** of the four PCs:

```
copy <repository>\station-agent C:\station-agent /r /e

cd C:\station-agent
py -3.10 -m venv .venv
.\.venv\Scripts\python -m pip install -r requirements.txt
```

Dependencies are a web framework, a server, a configuration loader and **process-inspection**. The process library matters:



Install the process library even though the service starts without it

Without it the agent cannot detect workstation processes, so the multi-station board never leaves "awaiting agent", and per-station CPU, memory and disk telemetry is absent. The service still runs, so the failure is silent.



Optional GPU telemetry

To show GPU cards on the PC monitoring screen, install the GPU bindings as well — they are not in the requirements file:

```
.\.venv\Scripts\python -m pip install pynvml
```

This needs a vendor GPU driver. Without it the card falls back to a browser-reported approximation.

11.3 Configure

```
copy .env.example .env
```

```
# The NAS root AS MAPPED ON THIS PC (drive letter or UNC path).
WATCH_ROOT=X:\nas\360_images

# Overrides, only if your layout differs from Section 9.
# OUT_STAGE=03_Stitching
# IN_STAGE=02_Blurring

# 1 = count capture POINTS instead of stitched images.
# Defaults to 1 for blur, 0 for the others.
# POINT_MODE=0

# Substrings matched against running process names/paths, comma separated.
# Match your actual software.
PROCESS_NAMES=ptgui.exe,ptgui.exe

# Watch loop period, and the "file is still growing" window in seconds.
SCAN_INTERVAL_SEC=5
GROWING_WINDOW_SEC=90

# Shared secret. Required unless the agent is unreachable from anywhere
# but this LAN.
AGENT_TOKEN=<generate a long random string per station>

# Bind address and port. Keep 0.0.0.0 only with a firewall rule.
AGENT_HOST=0.0.0.0
AGENT_PORT=8000
```

Use a **different token per station**, and register each one in `STATION_AGENT_TOKENS` (Section 8.4).

11.4 Verify

```
.\.venv\Scripts\python app.py
```

In a second shell:

```
curl http://localhost:8000/health
```

```
{
  "status": "ok", "agent_version": "1.0.0", "station_id": "stitch",
  "hostname": "PC2", "uptime_sec": 86400,
  "cpu_usage": 12, "cpu_cores": 16, "ram_total_gb": 32, "ram_used_gb": 14.2
}
```

Then check the station payload:

```
curl -H "Authorization: Bearer <AGENT_TOKEN>" http://localhost:8000/api/station
```

Confirm two things:

- the watch error is **null** — meaning `WATCH_ROOT` resolved on this PC, and
- the output section lists your subgrids.

If a watch error is present, the path is wrong for this machine. A drive letter that exists on one workstation frequently does not exist on another.

The agent binds port 8000 with no transport encryption of its own. Restrict it to the local network:

```
New-NetFirewallRule -DisplayName "GeoSphere Station Agent" `
  -Direction Inbound -LocalPort 8000 -Protocol TCP `
  -Action Allow -Profile Domain
```

The Cloudflare Tunnel is outbound-only and needs no inbound rule.

11.6 Autostart

Windows — Task Scheduler → Create Task:

- Trigger: **At startup**
- Action: **Start a program** → C:\station-agent\run_agent.bat
- Start in: C:\station-agent

Linux — a systemd unit pointing at the virtual environment's server binary.

Configuring the application itself, including one step that currently has no screen.

12.1 Create your first project

A project is **mandatory** — every query in the platform is scoped to the active project, and the dashboard forces project selection immediately after login. The wizard's only required field is the campaign name.

12.2 Required settings

Open **Settings** → **Project and Map Settings**. The form is read-only unless you are an Administrator.

SETTING	WHY IT IS REQUIRED
Supabase REST endpoint URL	No database client without it
Public anon API key	No sign-in without it
Database host provider	Selects your provider topology
Storage provider	Selects the panorama URL builder
Provider bucket / domain / container	Without it, publishing stays disabled
Default basemap provider	Map backdrop

Storage settings vary by provider — a bucket name, a CDN domain, or a bucket and region pair.



Build-time variables are not set here

The Supabase and map variables are **build-time** settings configured on Cloudflare Pages (Section 8.3). This screen only pre-fills the form; it does not set them. Changing them here has no effect on the running application.

12.3 Two settings in an unexpected place

The working directory and worker endpoint are **not** on the admin screen. Set them in **NAS and Daemon** → **Capacity and Volumes**:

SETTING	PURPOSE
NAS working base path	Root that deliverable paths resolve against
Worker URL	Endpoint used when not proxying

12.4 Station address configuration



No database step required

Station addresses are set in the interface: **Settings** → **Project & Map Settings** → **4. Production Pipeline & Workstations**. Enter the address for each of the four processing PCs and press **Save**.

Earlier revisions had no screen for this and required a settings-table SQL update, tracked as limitation 14.1. That defect is fixed. If you are upgrading, no migration is needed — fill in the addresses and save.

Agent port	Station agent listen port	8080
Live desktop channel	RDP (mstsc) or noVNC (browser)	RDP
RDP port	Port for the one-click mstsc launch	3389
noVNC port	websockify port for the in-browser desktop pane	blank
noVNC public URL	HTTPS noVNC endpoint; accepts {ip} and {port}	blank
Agent token	Shared secret matching the PC's AGENT_TOKEN	blank
Enabled	Whether the station is part of the active pipeline	on

The section header shows **N of 4 stations configured**. A station with no address reports itself as **Unconfigured** and the rest of the board keeps working, so PCs can be brought online one at a time. A mistyped address is flagged *Check address* rather than blocked, so a hostname that does not resolve to a dotted IP will not stop you saving.

The live desktop pane needs a public URL on HTTPS

A browser blocks a private `http://` VNC iframe as mixed content, so on an HTTPS deployment the pane requires the public noVNC URL above. Leave it blank and the pane stays hidden for that station; RDP launch still works from the office LAN.

Also sign out and back in after saving. The dashboard caches project settings in browser storage, so a hard refresh or a fresh session is needed to pick up the change. The stored settings type now declares the address and port keys, so exporting and re-importing project settings no longer drops them — a defect that affected earlier builds.

Topology is not editable here. Station order, software labels and the NAS stage folder templates come from the shipped four-station topology so the pipeline contract cannot drift per site. Only the address and remote-access fields are per-site.

In **proxy** mode — the default — the agents are reached through the Pages Functions using the hostname map, so the addresses matter less. Set them anyway so that local network fallback and the monitoring views work.

12.5 Roles and capabilities

Grant capabilities per operator in **Administration** → **Roles**. Remember that only eight capabilities are enforced by the database (Section 6.3) — do not rely on the others for compliance.

Two acceptance tiers, and the sign-off record.




Figure 8. Verification tiers. Tier B depends on Tier A plus the on-premises half; the sign-off sheet is in the appendix.

13.1 Tier A — dashboard verified

About thirty minutes. Requires no NAS. Capture a screenshot at each pass signal for the acceptance record.

#	ACTION	PASS SIGNAL
1	Open the deployed URL	Landing page renders with the globe animation
2	Sign in as the seeded Administrator	No error banner
3	Project picker appears	Zero projects forces the wizard
4	Complete the wizard — campaign name is the only required field	Loading indicator, then the dashboard
5	Expand the left navigation rail	All workspaces listed; Settings and Administration unlocked
6	Settings → Diagnostics → Ping	Connected , database operational , storage operational , a real latency figure
7	Expand Runtime Environment	Database and map variables show values, not "not set"
8	Administration → System Health → Run Diagnostics	Database operational with a real millisecond figure
9	Administration → Users	Your row present with role Administrator
10	Administration → Roles	The capability matrix renders
11	Settings → Test database connection and latency	A number replaces "not measured"
12	Settings → set the provider bucket → Verify manifest	Manifest OK with a frame count
13	Save settings	Success confirmation, and an entry in the audit log
14	Data Management → import a two-row CSV with <code>filename</code> , <code>latitude</code> , <code>longitude</code>	Columns auto-map; duplicate check reports 0
15	Road Analysis → Import Data → add a small GeoJSON	Feature count renders; transport badge reads worker
16	Dashboard	The embedded WebGIS renders a basemap
17	Header → audit log icon	The step-13 entry appears



Expected cosmetic artefact — not a fault

The diagnostics panel reports **WebGIS: unknown** and **Realtime: unknown** even on a perfectly healthy installation, because those probes are not implemented. Ignore them and judge health from the database, storage and latency

Cloudflare Pages and rebuild (Section 8.3).

13.2 Tier B — production pipeline verified

Tier A cannot prove the on-premises half. Continue once the worker, agents and tunnels are running.

#	CHECK	PASS SIGNAL
1	Scan for subgrids through the tunnel	Your subgrid is listed
2	NAS and Daemon workspace	Folder browser lists the stitching stage; storage overview shows the volume
3	Production Hub → flight board	All four cards show an agent, not "awaiting agent"
4	PC Monitoring	Live processor, memory and disk per station
5	Production Hub → intake and pairing	Subgrid detected, survey folders listed, telemetry rows paired to imagery
6	Quality inspection station	Frames step through the 360 viewer; a defect flag writes to the ledger
7	Cloud bucket gate	Inventory verified and gate approved
8	WebGIS release gate → publish	Rows land in the published table
9	Main Dashboard	The published trajectory renders

 **Steps 3 and 4 depend on Section 12.4**
Without the station registration those cards cannot reach the agents, however correctly the tunnels are configured.

 **Bucket gate limitation**
Inventory verification works only with Supabase object storage. Other providers cannot complete that gate today — see item 14.3.

13.3 Final smoke test

Run the repository's smoke checklist, which adds the authentication boundary, image revalidation and privilege checks. Record the outcome in the appendix sheet.

Product defects that affect a client installation. Read before go-live.

14.1 Station addresses must be set before the board is useful FIXED

Impact. Until each station has an address saved, the flight board reports that station as "Unconfigured", the rename action is unavailable for it, and bucket upload for non-Supabase providers cannot run for it. A station without an address does not affect the rest of the board, so PCs can be brought online one at a time.

Resolution. Earlier revisions had no screen for this and required a settings-table SQL update, and a settings export silently dropped the address and port keys. Both are fixed: set the addresses in **Settings** → **Project & Map Settings** → **4. Production Pipeline & Workstations** (Section 12.4), and the stored settings type now declares those keys. No migration is required when upgrading.

14.2 Bucket upload works for only two providers

Impact. When a bucket upload includes the manifest, the agent builds its command incorrectly for five of the seven supported providers, and the upload fails with an operating system error.

PROVIDER	UPLOAD WORKS
Cloudflare R2	Yes
Network-share copy	Yes
AWS S3, Google Cloud Storage, Azure, Wasabi, Supabase CLI	No — upload without the manifest option

Workaround. Use the R2 or network-share path, or run the provider command-line tool yourself and skip the in-app upload.

14.3 Bucket gate verification is Supabase-only

Impact. The inventory-verification step of the cloud bucket gate refuses every provider except Supabase object storage. R2, S3 and Azure deployments **cannot pass the gate**.

Workaround. Use Supabase object storage as the panorama bucket, or accept that the gate cannot be completed and publish through the release gate only.

14.4 The in-app tour and help text are out of date

Impact. The first-run interactive tour and the help dialog still describe the pre-refactor interface. They will mislead operators and new staff.

Workaround. Dismiss the tour and train against this guide instead.

14.5 Default session timeout is thirty minutes

Impact. Users are signed out after thirty minutes of inactivity. Activity inside embedded map and panorama frames is invisible to the dashboard, so someone actively working a 360 viewer can be signed out mid-task.

Workaround. Raise the session timeout in Settings if this is disruptive.

Impact. The delete capability is granted at the database level to both Administrator and Survey Operator, so the approval queue can be bypassed with a direct API call.

Workaround. Do not grant the Survey Operator role where your process requires approval.

14.6b Disable revokes permissions, but not the sign-in account

Impact. **Disable** in Administration → Users sets the directory row to **Disabled**. Since migration **0030** that resolves to read-only at the database level, so the control genuinely revokes.

It does **not** remove the person's Supabase sign-in account. They can still sign in; they simply see nothing they are not permitted to see. There is deliberately no delete control in the dashboard — deleting the row would leave no record of the revocation, and role resolution would fall back to the token claim, restoring the elevated role.

Workaround. To stop someone signing in at all, remove or disable the user in **Supabase → Authentication → Users**.

14.7 Retired configuration names

NAME	STATUS
VITE_NAS_API_ENABLED	Never existed in code. Documented in older notes
CONCURRENCY , MAX_ACTIVE_JOBS , MAX_QUEUE_DEPTH , WORKER_JOB_DB	Retired with the job queue. Setting them does nothing
MASK_NEIGHBOR_WARP , LAZY_MODEL_DOWNLOAD	Retired. Not read anywhere
Worker job, release and metrics endpoints	Removed. A "not found" response is expected, not a fault
Worker port 8787	Never correct. Use 8000
0019_map_shares_auth_only.sql	No such file. The real file is 0019_map_shares_public_access.sql

14.8 Renaming requires a panoramas folder

Impact. The NAS rename action targets the **panoramas/** subfolder of a survey run. If it does not exist the request is rejected.

Workaround. Use the tiled stitched layout (Section 9.1), or rename directly on the filesystem.

Symptom to cause. Entries marked silent fail without an error message.



This platform prefers to report "unavailable" over raising an error

Several failures are therefore silent. When something is not working, check the table below before assuming a data problem.

SYMPTOM	LIKELY CAUSE	FIX
Blank frame where the map should be (silent)	Map variable unset	Set it on Pages, then rebuild (8.3)
"Worker proxy is not configured" (silent)	Proxy variables absent in the development file	Add both, restart the dev server (8.6)
Edge returns 503 on private routes	Runtime secrets missing	Set the API URL and worker token (8.4)
Edge returns 401 on a private route	Session missing or expired	Sign in again; confirm the server Supabase variables
Tunnel hostname returns 404	Ingress rule does not match the requested hostname	Check the ingress rule; restart the tunnel service
Tunnel hostname returns 401 and nothing else works	Correct — the token is required	Add the authorisation header; if it works the tunnel is fine (7.6)
Tunnel hostname keeps changing	A quick tunnel was used	Create a named tunnel with DNS routing (7.2)
Station cards show no-ip-configured	Address not saved for that PC	Set it in Settings → 4. Production Pipeline & Workstations (Section 12.4)
Station cards show unreachable	Tunnel, firewall, or agent not running	Check the agent is up, the ingress matches, port 8000 is allowed
Station cards show not-reporting	Agent answers but its watch loop is failing	Check the watch root on that PC (11.4)
Agent exits immediately at startup	Station identifier missing or invalid	Must be one of the four values in 11.1
No processor or memory figures on a station	Process library not installed	Install the requirements on that PC (11.2)
No GPU card on the monitoring screen	GPU bindings not installed	Install the optional GPU package (11.2)
Panoramas render broken (silent)	Bucket is private, or the wrong bucket name	Make it public; re-check the setting (4.1)
"No resolvable panorama URL" in quality review	Storage provider or bucket incomplete	Re-check Section 12.2
Subgrid scan returns an empty list (silent)	Folder tree mismatch, or base path wrong	Re-check Sections 9.1 and 10.2
Coordinates not paired to imagery (silent)	Run folder names differ between the two stage folders	Make them identical (9.1)
Publish button permanently disabled	Provider bucket not configured	Set the bucket in Settings (12.2)
Cannot reach the Administration workspace	Seeded role resolved to Viewer	Verify the role query (5.2)

Migration error mentioning a security function	Applied out of order	Re-apply in order (3.3)
Summary view missing after migrations	PostGIS was not enabled before the first migration	Enable PostGIS, re-run the view block (3.2)
Diagnostics show WebGIS or Realtime unknown	Expected — probe not implemented	Ignore (13.1)
Database changes not appearing after an edit	Settings cached in the browser	Sign out and back in, or hard refresh

Lookup tables. Precedence is project settings, then environment variables, then built-in defaults.

16.1 Browser build-time variables

VARIABLE	REQUIRED	PURPOSE
VITE_SUPABASE_URL	Yes	Project base URL
VITE_SUPABASE_ANON_KEY	Yes	Public anon / publishable key
VITE_MAP_URL	Yes	WebGIS application base. No fallback
VITE_SUPABASE_KEY	No	Legacy alias for the anon key
VITE_WORKER_API_MODE	No	proxy (default) or direct
VITE_STATION_AGENT_MODE	No	proxy in production, direct in development
VITE_SUPABASE_BUCKET / VITE_STORAGE_BUCKET	No	Panorama bucket. Default MMS_PIC
VITE_STORAGE_PROVIDER	No	Storage provider selection
VITE_R2_BUCKET / VITE_R2_DOMAIN	No	Cloudflare R2 target
VITE_IMAGE_CDN_URL	No	Generic image CDN base
VITE_S3_BUCKET / VITE_S3_REGION	No	S3 target. Region default ap-southeast-1
VITE_GCS_BUCKET	No	Google Cloud Storage target
VITE_AZURE_ACCOUNT / VITE_AZURE_CONTAINER	No	Azure target
VITE_WASABI_BUCKET / VITE_WASABI_REGION	No	Wasabi target. Region default us-east-1
VITE_NAS_SERVER_URL	No	Direct NAS image origin
VITE_NAS_WORK_BASE_PATH	No	Default NAS working directory
VITE_PRODUCTION_API_URL	No	Worker base URL — only when mode is direct
VITE_PRODUCTION_API_KEY	No	Static worker secret — direct mode only. Prefer proxying
VITE_ROAD_GEOMETRY_BUCKET	No	Private layer geometry bucket
VITE_ROAD_EXTRACTION_ROUTE	No	overpass (default) or custom
VITE_ROAD_EXTRACTION_URL	No	Road-network endpoint, browser fallback only
VITE_ROAD_EXTRACTION_PROXY	No	Server-side proxy — preferred
VITE_ROAD_EXTRACTION_DIRECT	No	Bypass the proxy for debugging
VITE_MAPLIBRE_WORKER_URL	No	Self-hosted map worker for strict content policies
VITE_SHARE_LIVE_MODE	No	Embed the live map in share pages

VITE_... (right variables)	REQUIRED	PURPOSE
VITE_SENTRY_DSN	No	Error reporting. Omit to disable
VITE_DATA_QUIET	No	Suppress non-fatal analysis console noise

16.2 Edge runtime variables

Server-side only — see Section 8.4 for the full table with requirements and purposes.

16.3 NAS worker variables

VARIABLE	REQUIRED	DEFAULT
NAS_BASE_PATH	Yes	Survey tree root. Paths outside it are rejected
NAS_WORKER_TOKEN	Non-loopback	Blank means unauthenticated
NAS_WORKER_ID	No	Machine hostname
STORAGE_CACHE_TTL	No	30 seconds
NAS_LOG_JSON	No	0

16.4 Station agent variables

VARIABLE	REQUIRED	DEFAULT
STATION_ID	Yes	None — the service exits
WATCH_ROOT	No	Worker base path variable, else blank
OUT_STAGE	No	Per station (11.1)
IN_STAGE	No	Per station (11.1)
POINT_MODE	No	1 for blur, else 0
SUBGRID_PATTERN	No	^[A-Za-z]{1,3}[0-9A-Za-z]{2,11}\$
PROCESS_NAMES	No	Empty — no process detection
SCAN_INTERVAL_SEC	No	5
GROWING_WINDOW_SEC	No	90
AGENT_TOKEN	No	Blank means unauthenticated
AGENT_HOST / AGENT_PORT	No	0.0.0.0 / 8000

16.5 NAS worker HTTP surface

METHOD	ROUTE	AUTH	PURPOSE
--------	-------	------	---------

GET	/api/tables/{path}	Yes	Directory listing with image counts
GET	/api/storage	Yes	Volume usage, per top-level folder
GET	/api/nas-scan?action=	Yes	Survey scans
GET	/api/images/{path}	Yes	Panorama preview with conditional requests

Scan actions: `subgrids`, `survey-folders`, `read-csv`, `folder-images`, `final-images`, `registry`. An unknown action returns **400**.

16.6 Station agent HTTP surface

METHOD	ROUTE	PURPOSE
GET	/health	Liveness and host telemetry
GET	/api/station	Full station snapshot: task, output, capture points
GET	/api/station/{id}	Snapshot, only for the agent's own identifier
POST	/api/rename	Rename files within a stage folder
POST	/api/sync-bucket	Start a bucket upload
GET	/api/sync-bucket/{job_id}	Poll an upload

16.7 Capabilities enforced by the database

`runQaqc`, `reviewQaqc`, `viewAll`, `manageDatasets`, `manageSettings`, `manageUsers`, `approveDeletions`, `deleteData`.

The other thirteen capabilities in the interface matrix are presentation-only. See Section 6.3.

16.8 Continuous integration gates

GATE	RUNS
Build	Type check, then the production bundle
Test	Unit and integration suite
Lint	Static analysis — blocking on errors
Python	Compile and unit tests for the worker, on two interpreter versions

There is **no deployment automation** in the repository — the edge platform's Git integration is the deployment mechanism. The station agents and the edge functions are not covered by the automated gates.

Complete and sign once the installation is accepted.

A.1 Installation details

Client

Site / location

Supabase project URL

Web application URL

NAS worker hostname

NAS base path

Station agent hostnames

Guide revision applied

Installation date

A.2 Acceptance checklist

DONE	ITEM	VERIFIED BY	DATE
☐	PostGIS enabled before the first migration		
☐	All 29 migrations applied in order (0011 is a test, run separately)		
☐	Security test script passed (6.2 Tier 1)		
☐	Privilege boundary proven (6.2 Tier 2)		
☐	Panorama bucket created and public		
☐	First administrator bootstrapped and role verified		
☐	Operator accounts created with roles assigned		
☐	NAS folder tree matches Section 9		
☐	NAS worker responding with a valid token		
☐	Four station agents reporting		

☐	Five tunnels created and routed over HTTPS
☐	Tunnels and services set to start automatically
☐	Edge project root directory set to the repository root
☐	Build-time and runtime variables configured
☐	Station addresses registered (12.4)
☐	Tier A verification complete (13.1)
☐	Tier B verification complete (13.2)
☐	Smoke test complete (13.3)
☐	Known limitations in Section 14 reviewed with the client

A.3 Signatures

ROLE	NAME	SIGNATURE AND DATE
Client — IT / infrastructure		
Client — survey operations		
Vendor — delivery engineer		



Source documents

This guide consolidates the setup documentation held in the platform repository under `docs/Production Setup/` , together with the worker HTTP contract and the smoke-test checklist. Where this guide and an older note disagree, **this guide is current.**